

ALEX CHEN

Chicago, IL | alex.chen.ds@example.com | (312) 555-0198 | xx | xx

SUMMARY

Data Scientist with 5 years of experience building machine learning models, experimentation frameworks, and analytics pipelines across e-commerce, fintech, and mobility. Skilled in Python, SQL, machine learning, A/B testing, causal inference, cloud data workflows, and translating ambiguous business problems into production-ready data products.

TECHNICAL SKILLS

Languages: Python, SQL, R **ML/Stats:** XGBoost, LightGBM, Random Forest, Logistic Regression, NLP, Time Series, Clustering, A/B Testing, Causal Inference

Data/Cloud: pandas, NumPy, scikit-learn, PySpark, Airflow, dbt, Snowflake, BigQuery, AWS, PostgreSQL

Visualization/MLOps: Tableau, Looker, Power BI, MLflow, Docker, Git, CI/CD, model monitoring

EXPERIENCE

Data Scientist II | BrightCart Technologies - Chicago, IL | Jul 2023 - Present

- Built an XGBoost churn prediction model using behavioral features from 12M+ user sessions, increasing retention campaign reactivation rate by 18%.
- Developed personalized product ranking models from purchase history, search behavior, and item embeddings, improving click-through rate by 11% in A/B tests.
- Designed and analyzed pricing, search, and checkout experiments that supported product decisions tied to an estimated \$4.2M annualized revenue lift.
- Created automated Looker and MLflow monitoring dashboards to track feature drift, model stability, and KPI degradation; mentored 2 junior analysts.

Data Scientist | NorthBridge Financial Analytics - New York, NY | Jun 2021 - Jun 2023

- Developed credit default risk models with logistic regression, Random Forest, and LightGBM on 3M+ loan applications, improving approval precision by 14%.
- Built a PySpark fraud detection pipeline to identify anomalous application patterns, reducing manual review volume by 22%.
- Implemented Airflow and Snowflake feature workflows standardizing 150+ borrower, transaction, and bureau-derived features across modeling teams.
- Conducted fairness and bias analyses across demographic proxy groups and recommended threshold adjustments to reduce performance gaps.

Junior Data Scientist | UrbanPulse Mobility - Boston, MA | Aug 2019 - May 2021

- Forecasted hourly mobility demand using time series models, weather, calendar, and local event signals, reducing forecast error by 17%.
- Built clustering models to segment demand zones and identify high-utilization neighborhoods for fleet rebalancing and expansion planning.
- Automated SQL data quality checks and weekly Python reporting workflows, reducing duplicated-trip reporting errors and saving 8 analyst hours per week.

SELECTED PROJECTS

Real-Time Customer Churn Scoring: Built a FastAPI, Docker, and MLflow prototype serving daily churn probabilities with SHAP-based driver explanations for marketing teams.

Job Recommendation Engine: Ranked job descriptions against candidate profiles using embeddings, structured skill extraction, and filters for years of experience, location, sponsorship, and function.

Financial Transaction Anomaly Detection: Used Isolation Forest and robust statistical features to flag unusual transaction amount, velocity, merchant-category, and user-history patterns.

EDUCATION & CERTIFICATIONS

Northeastern University - B.S. in Data Science and Economics, Boston, MA, May 2019

Certifications: AWS Certified Machine Learning - Specialty; Google Cloud Professional Data Engineer; Tableau Desktop Specialist